

Laplacian relaxation optimization

DIETZ T., JABER A., DONNENFELD T.

04 November 2025

1 Algorithmic Changes

1.1 Matrix Update

The original implementation copied the new array into the old one element by element:

```
1 void laplace_copy(double *A, double *Anew, int n, int m) {
2     for (int j = 0; j < n; j++) {
3         for (int i = 0; i < m; i++) {
4             A[j*m + i] = Anew[j*m + i];
5         }
6     }
7 }
```

Listing 1: initial_code.c – full matrix copy

A better version is to simply swap the pointers to avoid the whole copy process. We just have to be careful with the borders of the uninitialized matrix, which by default have not been set to their constant values.

```
1 double *Atmp = A;
2 A = Anew;
3 Anew = Atmp;
```

Listing 2: final_code.c – pointer swap

This change prevents an entire pass over the matrix for each iteration.

1.2 Error Computation

In the initial implementation, the error was computed in a separate function, requiring a full matrix reading after each iteration.

```
1 double laplace_error(double *A, double *Anew, int n, int m) {
2     double error = 0.0;
3     for (int j = 1; j < n - 1; j++) {
4         for (int i = 1; i < m - 1; i++) {
5             error = fmax(error, fabs(Anew[j*m + i] - A[j*m + i]));
6         }
7     }
8     return error;
9 }
```

Listing 3: initial_code.c – separate error computation

In the optimized version, the error is computed directly during the relaxation step, removing the need for an extra matrix pass:

```

1 #pragma omp parallel for reduction(max:error) collapse(2)
2 for (int j = 1; j < n - 1; j++) {
3     for (int i = 1; i < m - 1; i++) {
4         Anew[j*m + i] = 0.25f * (A[j*m + i + 1] + A[j*m + i - 1] +
5                                     A[(j-1)*m + i] + A[(j+1)*m + i]);
6         error = fmaxf(error, fabsf(Anew[j*m + i] - A[j*m + i]));
7     }
8 }
```

Listing 4: final_code.c – integrated error computation

2 Code-Level Optimizations

2.1 Matrix index access swapping

Changing the way data from the matrices is accessed so that less memory fetches are made (Cache is reused since fetches bring contiguous blocks of memory). This makes a big difference because the first scenario generates the most possible amount of cache misses, and reordering the loops generates the most hits.

```

1 for ( i=1; i < m-1; i++ )
2     for ( j=1; j < n-1; j++ )
3         out[j*m+i]= stencil(in[j*m+i+1], in[j*m+i-1], in[(j-1)*m+i], in[(j+1)*m+i]);
```

Listing 5: initial_code.c

```

1 for ( j=1; j < n-1; j++ ) {
2     for ( i=1; i < m-1; i++ ) { // i moves in the inner loop
3         // ...
4 }
```

Listing 6: final_code.c

2.2 Parallelization with OpenMP

The optimized version introduces OpenMP directives to parallelize both the computation and the error reduction step:

```

1 #pragma omp parallel for reduction(max:error) collapse(2)
2 for (int j = 1; j < n - 1; j++) {
3     for (int i = 1; i < m - 1; i++) {
4         // Laplace update + error calculation
5     }
6 }
```

Listing 7: final_code.c – OpenMP parallel loop

2.3 Division Replacement

Replacing divisions with multiplications improves performance by avoiding expensive floating-point division operations:

```

1 Anew[j*m + i] = (A[j*m + i + 1] + A[j*m + i - 1] +
2                     A[(j-1)*m + i] + A[(j+1)*m + i]) / 4.0;
```

Listing 8: initial_code.c

```

1 Anew[j*m + i] = (A[j*m + i + 1] + A[j*m + i - 1] +
2                     A[(j-1)*m + i] + A[(j+1)*m + i]) * 0.25f;
```

Listing 9: final_code.c

2.4 Index Precomputation

Avoiding repeated multiplications inside the inner loop reduces instruction count:

```

1 int idx = j * m + i;
2 Anew[idx] = 0.25f * (A[idx + 1] + A[idx - 1] +
3                         A[idx - m] + A[idx + m]);

```

Listing 10: final_code.c – precomputed indices

3 Conclusion

The initial execution time was about 95 seconds. Through these different optimizations, while keeping the same initial problem values, we managed to get it down to 0.61 seconds which is 157 times faster.

Extra things to notice :

- CPU used went from 1 to 5.8
- Instruction throughput went from around 650M/s to 14G/s

```

Jacobi relaxation Calculation: 4096 x 4096 mesh, maximum of 100 iterations
 10, 0.246094
20, 0.176197
30, 0.144464
40, 0.125371
50, 0.112275
60, 0.102578
70, 0.095025
80, 0.088928
90, 0.083871
100, 0.079589
Total Iterations: 100, ERROR: 0.079589, A[32][32]= 0.006335

Performance counter stats for 'LAP2':
      95.661,45 msec task-clock:          #   1,000 CPUs utilized
          0 context-switches:u        #   0,000 /sec
          0 cpu-migrations:u       #   0,000 /sec
         1.159 page-faults:u        # 12,116 /sec
417.375.116.140  cpu_core/cycles:u/    #   4,363 G/sec
62.023.889.209  cpu_core/instructions:u/ # 648,369 M/sec
15.086.652.076  cpu_core/branches:u/   # 157,709 M/sec
     1.308.047  cpu_core/branch-misses:u/ # 13,674 K/sec

95,670057061 seconds time elapsed

95,588438000 seconds user
 0,013929000 seconds sys

```

Figure 1: Initial version runtime

```

Jacobi relaxation Calculation: 4096 x 4096 mesh, maximum of 100 iterations
 10, 0.246094
20, 0.176197
30, 0.144464
40, 0.125371
50, 0.112275
60, 0.102578
70, 0.095025
80, 0.088928
90, 0.083871
100, 0.079589
Total Iterations: 100, ERROR: 0.079589, A[32][32]= 0.006335

Performance counter stats for 'LAP2':
      3,595,30 msec task-clock:u          #   5,886 CPUs utilized
          0 context-switches:u        #   0,000 /sec
          0 cpu-migrations:u       #   0,000 /sec
         1.179 page-faults:u        # 327,928 /sec
14.255.285.703  cpu_core/cycles:u/    #   3,965 G/sec
50.318.479.976  cpu_core/instructions:u/ # 13,996 G/sec
11.742.747.893  cpu_core/branches:u/   #  3,266 G/sec
     394.286  cpu_core/branch-misses:u/ # 109,667 K/sec

0,610855730 seconds time elapsed

3,574385000 seconds user
 0,011965000 seconds sys

```

Figure 2: Optimized version runtime